


In [239]:

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
books = pd.read_csv("clean_books.csv")
```

In [240]:

```
1 sns.histplot(data=books , x='rating')
2 plt.show()
```



In [241]:

```
1 books.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 350 entries, 0 to 349
Data columns (total 5 columns):
#   Column  Non-Null Count  Dtype  
---  -
0    name    350 non-null     object  
1    author  350 non-null     object  
2    rating  350 non-null     float64  
3    year    350 non-null     int64   
4    genre   350 non-null     object  
dtypes: float64(1), int64(1), object(3)
memory usage: 13.8+ KB
```

In [242]:

```
1 books.value_counts("genre")
```

Out[242]:

```
genre
Non Fiction    179
Fiction        131
Childrens      40
dtype: int64
```

In [243]:

```
1 books['genre'].value_counts()
```

Out[243]:

```
Non Fiction    179
Fiction        131
Childrens      40
Name: genre, dtype: int64
```

In [244]:

```
1 books.value_counts("genre")
```


Out[244]:

```
genre
Non Fiction    179
Fiction        131
Childrens      40
dtype: int64
```

1

In [245]:

```
1 unemployment = pd.read_csv("clean_unemployment.csv")
2 sns.histplot(data=unemployment, x='2021', binwidth = 1)
3 plt.show()
```



In [246]:

```
1 books.dtypes
```

Out[246]:

```
name      object
author    object
rating     float64
year       int64
genre     object
dtype: object
```

In [247]:

```
1 books["year"] = books["year"].astype(int)
2 books.dtypes
```

Out[247]:

```
name      object
author    object
rating     float64
year       int64
genre      object
dtype: object
```

Validating categorical data

In [248]:

```
1 books["genre"].isin(["Fiction", "Non Fiction"])
```

Out[248]:

```
0      True
1      True
2      True
3      True
4     False
...
345    True
346    True
347    True
348    True
349    False
Name: genre, Length: 350, dtype: bool
```

In [249]:

```
1 books[books["genre"].isin(["Fiction", "Non Fiction"])].head()
```

Out[249]:

| | name | author | rating | year | genre |
|---|---|---------------------|--------|------|-------------|
| 0 | 10-Day Green Smoothie Cleanse | JJ Smith | 4.7 | 2016 | Non Fiction |
| 1 | 11/22/63: A Novel | Stephen King | 4.6 | 2011 | Fiction |
| 2 | 12 Rules for Life: An Antidote to Chaos | Jordan B. Peterson | 4.7 | 2018 | Non Fiction |
| 3 | 1984 (Signet Classics) | George Orwell | 4.7 | 2017 | Fiction |
| 5 | A Dance with Dragons (A Song of Ice and Fire) | George R. R. Martin | 4.4 | 2011 | Fiction |

In [250]:

```
1 books.select_dtypes("number").head()
```

Out[250]:

| | rating | year |
|---|--------|------|
| 0 | 4.7 | 2016 |
| 1 | 4.6 | 2011 |
| 2 | 4.7 | 2018 |
| 3 | 4.7 | 2017 |
| 4 | 4.8 | 2019 |

In [251]:

```
1 books["year"].min()
```

Out[251]:

2009

In [252]:


```
1 books["year"].max()
```

Out[252]:

2019

In [253]:

```
1 sns.boxplot(data=books , x = "year")  
2 plt.show()
```




In [254]:

```
1 sns.boxplot(data = books , x="year" , y ="genre")
```

Out[254]:

<AxesSubplot:xlabel='year', ylabel='genre'>



2

In [255]:

```
1 unemployment = pd.read_csv("clean_unemployment.csv")
2 not_oceania = ~unemployment['continent'].isin(["Oceania"])
3 not_oceania
```

Out[255]:

```
0      True
1      True
2      True
3      True
4      True
...
177    False
178     True
179     True
180     True
181     True
Name: continent, Length: 182, dtype: bool
```

3

In [256]:

```
1 unemployment["2021"].min()
2
```

Out[256]:

0.26

In [257]:

```
1 unemployment["2021"].max()
```

Out[257]:


33.56

In [258]:

```
1 sns.boxplot(data=unemployment, x="2021", y="continent")
```

Out[258]:

<AxesSubplot: xlabel='2021', ylabel='continent'>



Exploring groups of data

In [259]:

```
1 books.groupby("genre").mean()
```

Out[259]:

| | rating | year |
|-------------|----------|-------------|
| genre | | |
| Childrens | 4.780000 | 2015.075000 |
| Fiction | 4.570229 | 2013.022901 |
| Non Fiction | 4.598324 | 2013.513966 |

In [260]:

```
1 books.agg(["mean", "std"])
```

/var/folders/4j/bnvctt7152z6l5l6szd4m7wh0000gn/T/ipykernel_49443/1965544463.py:1: FutureWarning: ['name', 'author', 'genre'] did not aggregate successfully. If any error is raised this will raise in a future version of pandas. Drop these columns/ops to avoid this warning.
books.agg(["mean", "std"])

Out[260]:

| | rating | year |
|------|----------|-------------|
| mean | 4.608571 | 2013.508571 |
| std | 0.226941 | 3.284711 |

In [261]:

```
1 books.agg({"rating": ["mean", "std"], "year": ["median"]})
```

Out[261]:

| | rating | year |
|--------|----------|--------|
| mean | 4.608571 | NaN |
| std | 0.226941 | NaN |
| median | NaN | 2013.0 |

In [262]:


```
1 books.groupby("genre").agg(mean_rating=("rating", "mean"),
2 std_rating = ("rating", "std"),
3 median_year = ("year", "median")
4 )
```

Out[262]:

| | mean_rating | std_rating | median_year |
|-------------|-------------|------------|-------------|
| genre | | | |
| Childrens | 4.780000 | 0.122370 | 2015.0 |
| Fiction | 4.570229 | 0.281123 | 2013.0 |
| Non Fiction | 4.598324 | 0.179411 | 2013.0 |

In [263]:

```
1 sns.barplot(data=books, x = "genre" ,y = "rating")
2 plt.show()
```



4

In [264]:

```
1 unemployment.agg(["mean" , "std"])
```

/var/folders/4j/bnvctt7152z61516szd4m7wh0000gn/T/ipykernel_49443/2501974079.py:1: FutureWarning: ['country_code', 'country_name', 'continent'] did not aggregate successfully. If any error is raised this will raise in a future version of pandas. Drop these columns/ops to avoid this warning.
unemployment.agg(["mean" , "std"])

Out[264]:

| | 2010 | 2011 | 2012 | 2013 | 2014 | 2015 | 2016 | 2017 | 2018 | 2019 | 2020 | 2021 |
|------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|
| mean | 8.409286 | 8.315440 | 8.317967 | 8.344780 | 8.179670 | 8.058901 | 7.925879 | 7.668626 | 7.426429 | 7.243736 | 8.420934 | 8.390879 |
| std | 6.248887 | 6.266795 | 6.367270 | 6.416041 | 6.284241 | 6.161170 | 6.045439 | 5.902152 | 5.818915 | 5.696573 | 6.040915 | 6.067192 |

In [265]:

```
1 unemployment.groupby("continent").agg(["mean" , "std"])
```

/var/folders/4j/bnvctt7152z61516szd4m7wh0000gn/T/ipykernel_49443/564087925.py:1: FutureWarning: ['country_code', 'country_name'] did not aggregate successfully. If any error is raised this will raise in a future version of pandas. Drop these columns/ops to avoid this warning.
unemployment.groupby("continent").agg(["mean" , "std"])

Out[265]:

| | 2010 | | 2011 | | 2012 | | 2013 | | 2014 | | ... | 2017 | | 2018 | |
|---------------|-----------|----------|-----------|----------|-----------|----------|-----------|----------|-----------|----------|-----|----------|----------|----------|--|
| | mean | std | mean | std | mean | std | mean | std | mean | std | ... | mean | std | mean | |
| continent | | | | | | | | | | | | | | | |
| Africa | 9.343585 | 7.411259 | 9.369245 | 7.401556 | 9.240755 | 7.264542 | 9.132453 | 7.309285 | 9.121321 | 7.291359 | ... | 9.284528 | 7.407620 | 9.237925 | |
| Asia | 6.240638 | 5.146175 | 5.942128 | 4.779575 | 5.835319 | 4.756904 | 5.852128 | 4.668405 | 5.853191 | 4.681301 | ... | 6.171277 | 5.277201 | 6.090213 | |
| Europe | 11.008205 | 6.392063 | 10.947949 | 6.539538 | 11.325641 | 7.003527 | 11.466667 | 6.969209 | 10.971282 | 6.759765 | ... | 8.359744 | 5.177845 | 7.427436 | |
| North America | 8.663333 | 5.115805 | 8.563333 | 5.377041 | 8.448889 | 5.495819 | 8.840556 | 6.081829 | 8.512222 | 5.801927 | ... | 7.391111 | 5.326446 | 7.281111 | |
| Oceania | 3.622500 | 2.054721 | 3.647500 | 2.008466 | 4.103750 | 2.723118 | 3.980000 | 2.640119 | 3.976250 | 2.659205 | ... | 3.872500 | 2.492834 | 3.851250 | |
| South America | 6.870833 | 2.807058 | 6.518333 | 2.801577 | 6.410833 | 2.936508 | 6.335000 | 2.808780 | 6.347500 | 2.834332 | ... | 7.281667 | 3.398994 | 7.496667 | |

6 rows x 24 columns



5

In [266]:

```
1
2 continent_summary = unemployment.groupby("continent").agg(mean_rate_2021=("2021" , "mean"),
3                                                              std_rate = ("2021" , "std"))
4 continent_summary
5
```


Out[266]:

| | mean_rate_2021 | std_rate |
|---------------|----------------|----------|
| continent | | |
| Africa | 10.473585 | 8.131636 |
| Asia | 6.906170 | 5.414745 |
| Europe | 7.414872 | 3.947825 |
| North America | 9.155000 | 5.076482 |
| Oceania | 4.280000 | 2.671522 |
| South America | 9.924167 | 3.611624 |

6

In [267]:

```
1 sns.barplot(data=unemployment, x = "continent" ,y = "2021")
2 plt.show()
```



In [268]:

```
1 salaries = pd.read_csv("ds_salaries_clean.csv")
2
```

In [269]:

```
1 print(salaries.isna().sum())
```

```
Working_Year      0
Designation       0
Experience         0
Employment_Status 0
Employee_Location 0
Company_Size      0
Remote_Working_Ratio 0
Salary_USD        0
dtype: int64
```

In [270]:

```
1 threshold = len(salaries) * 0.05
2 print(threshold)
```

30.35

In [271]:

```
1 cols_to_drop = salaries.columns[salaries.isna().sum() <= threshold]
2 print(cols_to_drop)
```

```
Index(['Working_Year', 'Designation', 'Experience', 'Employment_Status',
      'Employee_Location', 'Company_Size', 'Remote_Working_Ratio',
      'Salary_USD'],
      dtype='object')
```

In [272]:

```
1 salaries.dropna(subset = cols_to_drop , inplace = True)
```

In [273]:

```
1 cols_with_missing_values = salaries.columns[salaries.isna().sum() > 0]
2 print(cols_with_missing_values)
```

Index([], dtype='object')

In [274]:

```
1 for col in cols_with_missing_values[:-1]:
2     salaries[col].fillna(salaries[col].mode()[0])
```

In [275]:

```
1 print(salaries.isna().sum())
```

```
Working_Year      0
Designation       0
Experience         0
Employment_Status 0
Employee_Location 0
Company_Size      0
Remote_Working_Ratio 0
Salary_USD        0
dtype: int64
```

In [276]:

```
1 salaries_dict = salaries.groupby("Experience")["Salary_USD"].median().to_dict()
2 print(salaries_dict)
```

```
{'Entry': 53948.0, 'Executive': 163694.5, 'Mid': 73465.0, 'Senior': 129380.0}
```

In [277]:

```
1 salaries["Salary_USD"] = salaries["Salary_USD"].fillna(salaries["Experience"].map(salaries_dict))
2 print(salaries.isna().sum())
```

```
Working_Year      0
Designation       0
Experience         0
Employment_Status 0
Employee_Location 0
Company_Size      0
Remote_Working_Ratio 0
Salary_USD        0
dtype: int64
```

7

In [278]:

```
1 planes = pd.read_csv("Airlines_unclean.csv", index_col = 0)
2 print(planes.isna().sum())
```

```
Airline      427
Date_of_Journey 322
Source       187
Destination   347
Route        256
Dep_Time     260
Arrival_Time 194
Duration     214
Total_Stops   212
Additional_Info 589
Price        616
dtype: int64
```

In [279]:

```
1 threshold = len(planes) * 0.05
2 print(threshold)
```

533.0

In [280]:

```
1 cols_to_drop = planes.columns[planes.isna().sum() <= threshold]
2 print(cols_to_drop)
```

```
Index(['Airline', 'Date_of_Journey', 'Source', 'Destination', 'Route',
       'Dep_Time', 'Arrival_Time', 'Duration', 'Total_Stops'],
      dtype='object')
```

In [281]:

```
1 planes.dropna(subset = cols_to_drop , inplace = True)
```

In [282]:

```
1 cols_with_missing_values = planes.columns[planes.isna().sum() > 0]
2 print(cols_with_missing_values)
```

```
Index(['Additional_Info', 'Price'], dtype='object')
```

In [283]:

```
1 for col in cols_with_missing_values[:-1]:
2     planes[col].fillna(planes[col].mode()[0])
```

In [284]:

```
1 print(planes.isna().sum())
```

```
Airline      0
Date_of_Journey  0
Source       0
Destination   0
Route        0
Dep_Time     0
Arrival_Time  0
Duration     0
Total_Stops   0
Additional_Info  300
Price        368
dtype: int64
```

Strategies for remaining missing data.

In [285]:

```

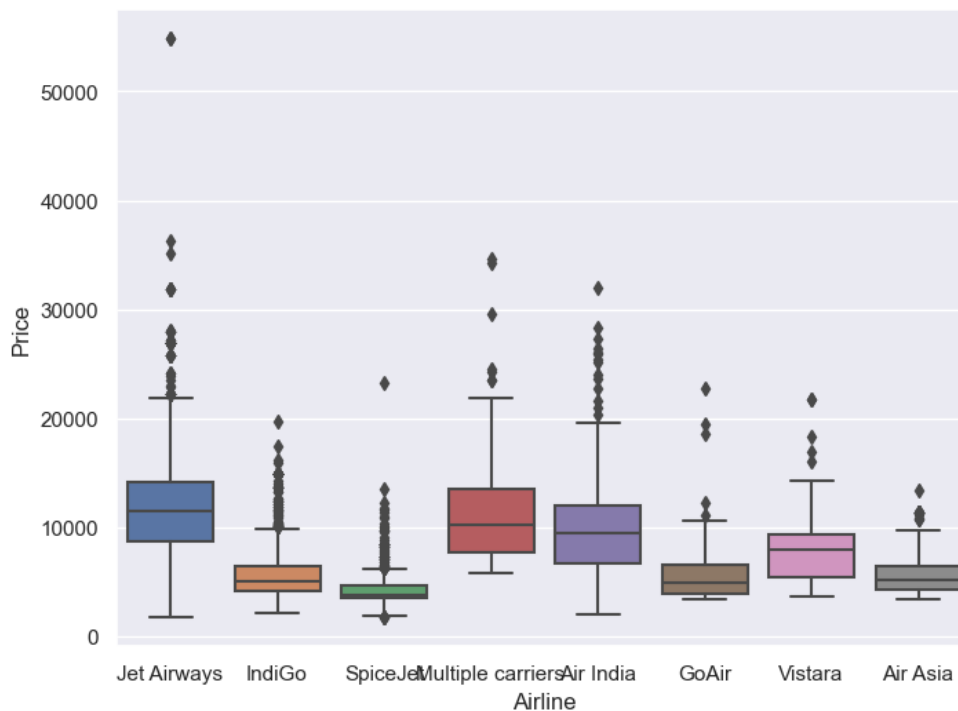
1 # Check the values of the Additional_Info column
2 print(planes["Additional_Info"].value_counts())
3
4 # Create a box plot of Price by Airline
5 sns.boxplot(data=planes, x='Airline', y='Price')
6 sns.set(rc={"figure.figsize":(8, 6)}) #width=8, #height=6
7 plt.show()
8

```

```

No info                6399
In-flight meal not included 1525
No check-in baggage included 258
1 Long layover         14
Change airports        7
No Info                2
Business class         1
Red-eye flight         1
2 Long layover         1
Name: Additional_Info, dtype: int64

```



8

In [286]:

```

1 planes = planes.drop(columns = ['Additional_Info'])
2 planes_dict = planes.groupby("Airline")["Price"].median().to_dict()
3
4
5 print(planes_dict)
6

```

```

{'Air Asia': 5192.0, 'Air India': 9443.0, 'GoAir': 5003.5, 'IndiGo': 5054.0, 'Jet Airways': 11507.0, 'Multiple carriers': 10197.0, 'SpiceJet': 3873.0, 'Vistara': 8028.0}

```

In [287]:

```

1 planes["Price"] = planes["Price"].fillna(planes["Airline"].map(planes_dict))
2 print(planes.isna().sum())

```

```

Airline                0
Date_of_Journey        0
Source                 0
Destination            0
Route                 0
Dep_Time              0
Arrival_Time          0
Duration              0
Total_Stops           0
Price                 0
dtype: int64

```

WORKING WITH CATEGORICAL DATA

In [288]:

```
1 salaries = pd.read_csv("ds_salaries_clean.csv")
2 print(salaries.select_dtypes("object").head())
```

```

      Designation Experience Employment_Status Employee_Location \
0      Data Scientist      Mid              FT                DE
1  Machine Learning Scientist    Senior              FT                JP
2      Big Data Engineer    Senior              FT                GB
3    Product Data Analyst      Mid              FT                HN
4  Machine Learning Engineer    Senior              FT                US

      Company_Size
0                L
1                S
2                M
3                S
4                L

```

In [289]:

```
1 print(salaries["Designation"].value_counts())
```

```

Data Scientist      143
Data Engineer      132
Data Analyst        97
Machine Learning Engineer  41
Research Scientist  16
Data Science Manager  12
Data Architect      11
Big Data Engineer   8
Machine Learning Scientist  8
Principal Data Scientist  7
AI Scientist        7
Data Science Consultant  7
Director of Data Science  7
Data Analytics Manager  7
ML Engineer         6
Computer Vision Engineer  6
BI Data Analyst     6
Lead Data Engineer  6
Data Engineering Manager  5
Business Data Analyst  5
Head of Data        5
Applied Data Scientist  5
Applied Machine Learning Scientist  4
Head of Data Science  4
Analytics Engineer  4
Data Analytics Engineer  4
Machine Learning Developer  3
Machine Learning Infrastructure Engineer  3
Lead Data Scientist  3
Computer Vision Software Engineer  3
Lead Data Analyst   3
Data Science Engineer  3
Principal Data Engineer  3
Principal Data Analyst  2
ETL Developer       2
Product Data Analyst  2
Director of Data Engineering  2
Financial Data Analyst  2
Cloud Data Engineer  2
Lead Machine Learning Engineer  1
NLP Engineer        1
Head of Machine Learning  1
3D Computer Vision Researcher  1
Data Specialist     1
Staff Data Scientist  1
Big Data Architect  1
Finance Data Analyst  1
Marketing Data Analyst  1
Machine Learning Manager  1
Data Analytics Lead  1
Name: Designation, dtype: int64

```

In [290]:

```
1 print(salaries["Designation"].nunique())
```

50

In [291]:

```
1 salaries["Designation"].str.contains("Scientist")
```

Out[291]:

```
0      True
1      True
2     False
3     False
4     False
...
602    False
603    False
604    False
605    False
606     True
```

Name: Designation, Length: 607, dtype: bool

In [292]:

```
1 salaries["Designation"].str.contains("Machine Learning|AI")
```

Out[292]:

```
0     False
1      True
2     False
3     False
4      True
...
602    False
603    False
604    False
605    False
606     True
```

Name: Designation, Length: 607, dtype: bool

In [293]:

```
1 job_categories = ["Data Science", "Data Analytics", "Data Engineering", "Machine Learning", "Managerial", "Consultant"]
```

In [294]:

```
1 data_scientist = "Data Scientist|NLP"
2 data_analyst = "Analyst|Analytics"
3 data_engineer = "Data Engineer|ETL|Architect|Infrastructure"
4 ml_engineer = "Machine Learning|ML|Big Data|AI"
5 manager = "Manager|Head|Director|Lead|Principal|Staff"
6 consultant = "Consultant|Freelance"
```

In [295]:

```
1 conditions = [  
2 (salaries["Designation"].str.contains(data_science)),  
3 (salaries["Designation"].str.contains(data_analyst)),  
4 (salaries["Designation"].str.contains(data_engineer)),  
5 (salaries["Designation"].str.contains(ml_engineer)),  
6 (salaries["Designation"].str.contains(manager)),  
7 (salaries["Designation"].str.contains(consultant))  
8 ]  
9 conditions
```

Out[295]:

```
[0      True  
1      False  
2      False  
3      False  
4      False  
...  
602     False  
603     False  
604     False  
605     False  
606     False  
Name: Designation, Length: 607, dtype: bool,  
0      False  
1      False  
2      False  
3      False  
4      False  
...  
602     False  
603     False  
604     False  
605     False  
606     False  
Name: Designation, Length: 607, dtype: bool,  
0      False  
1      False  
2      True  
3      False  
4      False  
...  
602     True  
603     True  
604     False  
605     False  
606     False  
Name: Designation, Length: 607, dtype: bool,  
0      False  
1      True  
2      True  
3      False  
4      True  
...  
602     False  
603     False  
604     False  
605     False  
606     True  
Name: Designation, Length: 607, dtype: bool,  
0      False  
1      False  
2      False  
3      False  
4      False  
...  
602     False  
603     False  
604     False  
605     False  
606     False  
Name: Designation, Length: 607, dtype: bool,  
0      False  
1      False  
2      False  
3      False  
4      False  
...  
602     False  
603     False  
604     False  
605     False  
606     False  
Name: Designation, Length: 607, dtype: bool]
```

In [296]:

```
1 import numpy as np
2 salaries["Job_Category"] = np.select(conditions, job_categories, default = "Other")
3 salaries["Job_Category"]
```

Out[296]:

```
0      Data Science
1      Machine Learning
2      Data Engineering
3      Other
4      Machine Learning
...
602    Data Engineering
603    Data Engineering
604      Other
605      Other
606    Machine Learning
Name: Job_Category, Length: 607, dtype: object
```

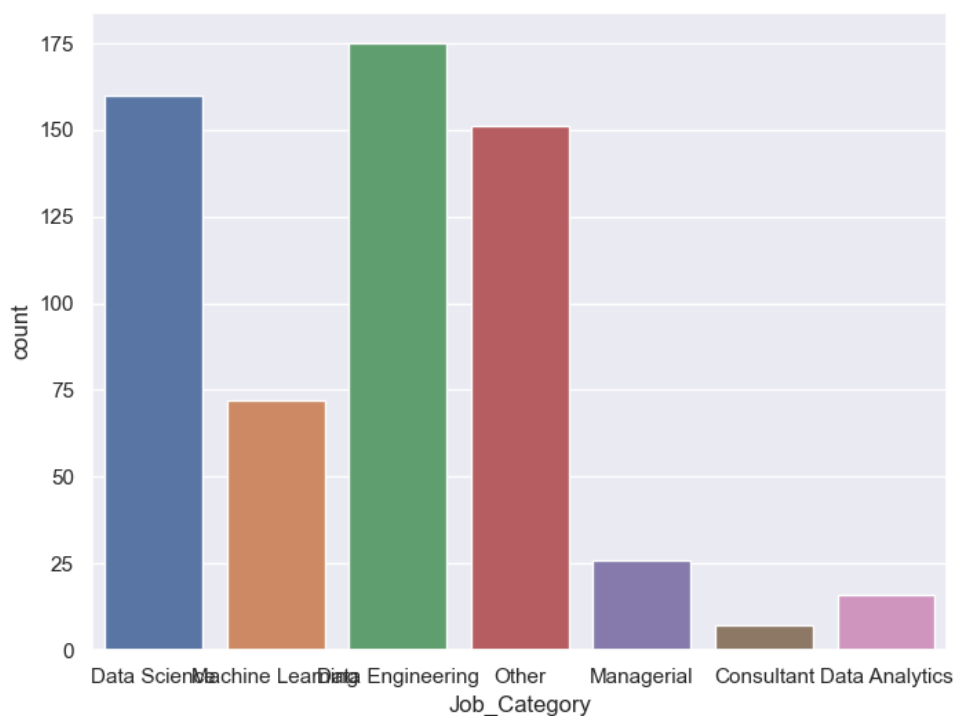
In [297]:

```
1 print(salaries[["Designation", "Job_Category"]].head())
```

| | Designation | Job_Category |
|---|----------------------------|------------------|
| 0 | Data Scientist | Data Science |
| 1 | Machine Learning Scientist | Machine Learning |
| 2 | Big Data Engineer | Data Engineering |
| 3 | Product Data Analyst | Other |
| 4 | Machine Learning Engineer | Machine Learning |

In [298]:

```
1 sns.countplot(data=salaries, x="Job_Category")
2 plt.show()
```



In [299]:

```
1 # planes = pd.read_csv("Airlines_unclean.csv")
2 # non_numeric = planes.select_dtypes("object")
3 # for col in non_numeric.columns:
4 #     print(f"Number of unique values in {col} column: ", non_numeric[col].nunique())
```

In [300]:

```
1 # planes["Duration"].head()
```


In [301]:

```
#1 import numpy as np
#2 # values = ["0h", "1h", "2h", "3h", "4h", "5h", "6h", "7h", "8h", "9h", "10h", "11h", "12h", "13h", "14h", "15h", "16h"]
3
4
5
6
#7 short_flights = "0h|1h|2h|3h|4h"
#8 medium_flights = "5h|6h|7h|8h|9h"
#9 long_flights = "10h|11h|12h|13h|14h|15h|16h"
10
#1 conditions = [
#2     (planes["Duration"].str.contains(short_flights)),
#3     (planes["Duration"].str.contains(medium_flights)),
#4     (planes["Duration"].str.contains(long_flights))
#5 ]
16
#7 values = [short_flights, medium_flights, long_flights]
18
19
#0 planes['Duration_Category'] = np.select(conditions, values, default='Other')
#1 planes['Duration_Category']
#2 # conditions
```

In []:

1

In []:

1

In []:

1